

Predictive Analytics for Bus Service Reliability

A Big Data Programming Project using PySpark

Project Title: Predictive Analytics for Bus Service Reliability: Classifying Breakdown Risk and Forecasting Delay Duration at Scale **Student Name & ID:** [Student Name] — [Student ID] **Module:** Big Data Programming Project **Supervisor:** Mr. Siddhartha Neupane **Date of Submission:** [Date]

1. Executive Summary

Urban bus networks generate large volumes of operational incident data that, if harnessed, can improve reliability and accountability. This project designs, implements and evaluates a genuine big-data predictive-analytics platform built on **Apache Spark (PySpark)**, using **1.29 million** real bus breakdown and delay records published by New York City Open Data. After distributed cleaning and feature engineering, **1,167,832** analysis-ready records were persisted to a Parquet store and a normalised SQLite star schema. Two supervised learning problems were addressed with Spark MLlib: a **class-imbalanced classification** of severe *breakdowns* versus minor *running-late* incidents, and a **regression** forecasting delay duration in minutes. Six models were compared; a Gradient-Boosted-Tree classifier achieved the best ROC-AUC of **0.847**, and a Random-Forest regressor achieved the lowest RMSE of **14.19 minutes** ($R^2 = 0.52$). The platform demonstrates partitioning, caching, broadcast joins, parameterised (injection-safe) SQL, and an interactive Streamlit dashboard, providing a transport authority with actionable, evidence-based insight into operator reliability and delay patterns.

2. Introduction

Problem statement. Which reported incidents are likely to be severe *breakdowns* rather than minor *delays*, and how long will a delay last? Answering these enables proactive resource deployment and operator benchmarking.

Purpose and scope. The goal is an end-to-end, reproducible pipeline — ingestion, storage, analytics, prediction and visualisation — that operates at genuine big-data scale and quantifies service reliability.

Relevance. In a smart-city context, converting raw telemetry into predictive insight reduces delays, informs infrastructure decisions and supports regulatory oversight — directly serving a **Transport Authority / Local Council** stakeholder.

Metric definitions (aligned to the brief, §4). We adopt three operational definitions: **Service Reliability** — the proportion of incidents resolved within an acceptable tolerance, operationalised as `on_time` (delay ≤ 15 min); **Travel-Time Variability** — the spread of delay duration, targeted directly by the regression model; and **Algorithmic Efficiency** — wall-clock runtime and Spark-stage utilisation, reported throughout. A composite **breakdown rate** (mean of `is_breakdown`) is defined as the operator-level severity indicator used for benchmarking.

Learning outcomes targeted. The project addresses **B1** (algorithmic complexity), **B2** (multi-language programming: PySpark, SQL, Python), **B4** (large-scale data science), **B6** (professional practice: Git, secure queries), **B7** (transferable communication skills) and **B8** (advanced work).

3. Literature Review

Predictive transport analytics is well established. Regression and tree-ensemble models are widely used for travel-time and delay prediction because they capture non-linear interactions between temporal, spatial and operator factors (Moreira-Matias et al., 2013). Gradient-boosted trees frequently outperform linear baselines on heterogeneous tabular transport data. For rare-event detection such as breakdowns, class imbalance is a recognised challenge, mitigated by cost-sensitive (weighted) learning and threshold-independent metrics such as ROC-AUC and PR-AUC (He & Garcia, 2009). Apache Spark has become the de-facto engine for distributed transport-data processing owing to its lazy DAG execution and in-memory caching (Zaharia et al., 2016). This project synthesises these strands: cost-sensitive ensembles on a distributed Spark backbone.

4. Data Collection & Preprocessing

Data source. The primary dataset is NYC Open Data's *Bus Breakdown and Delays* (id `ez4e-fazm`, **CC0 Public Domain**), the authoritative source behind the widely-used Kaggle mirror. It records every breakdown/delay logged by school- bus vendors (2015–2026): 21 attributes spanning temporal, operator, route, borough, reason and outcome fields.

Data scale. The ingested CSV contained **1,294,129 rows (~295 MB)** — over twelve times the 100,000-record threshold — satisfying the big-data requirement through an *extended date-range* strategy rather than synthetic augmentation.

Tools & justification (per stage). **PySpark** performed all large-scale work (quote-aware CSV parsing, transformation, joins, SQL, ML) because the data exceeds comfortable single-machine memory and benefits from partitioned parallelism. **Pandas** was used *only* at the final presentation boundary — converting small aggregated results for matplotlib/seaborn plotting and feeding the Streamlit layer — a deliberate, minimal memory-scale hand-off enabled by Apache Arrow.

Data cleaning. The free-text `How_Long_Delayed` field was parsed into numeric minutes with a regular-expression rule set (handling ranges, “MINS”, “HOUR”), then winsorised to a plausible 0–180-minute band. Timestamps were parsed with an explicit 12-hour pattern; **126,297 records (9.8 %)** carrying corrupt years (e.g. pre-1900 typos) were removed, leaving **1,167,832** clean rows. High-cardinality operator names were normalised (vendor codes stripped) and, for modelling, bucketed to the top-25 plus “OTHER”.

Merging strategy. The wide table was decomposed into a **star schema** — one `fact_incident` table linked by foreign keys to `dim_operator`, `dim_reason`, `dim_boro` and `dim_date` — and the dimensions were re-joined via **broadcast joins**, demonstrating relationship modelling across multiple related tables.

Challenges. A slow download link was overcome with gzip compression; a Spark 3.5 “ancient-datetime” Parquet-write error was resolved with a proleptic-Gregorian rebase configuration plus the

date-quality filter above.




Figure 1. Right-skewed distribution of parsed delay minutes (*skew* = 1.00, *median* = 30 min).

5. Methodology

Feature engineering. From timestamps we derived `hour`, `day_of_week`, `month`, `is_weekend`, a four-level `time_band` (AM/PM peak, midday, off-peak) and `season`. Categorical predictors (operator, borough, run type, reason, school-age) were `StringIndexer` → `OneHotEncoder` encoded and assembled with numeric features via `VectorAssembler`; a `StandardScaler` completed the shared pipeline.

Target design & leakage control. For the breakdown classifier, `reason` and `delay_minutes` were **deliberately excluded** — both are known only *after* the event — yielding an honest *a-priori* risk model. For delay regression, `reason` is legitimately available and is included.


Model selection. Three models per task were compared: **Logistic Regression**, **Random Forest**, **Gradient-Boosted Trees** (classification); **Linear**, **Random-Forest** and **GBT regressors** (regression). Class imbalance (9.6 % breakdowns) was addressed with inverse-frequency `classWeights`.

Data splitting & validation. An 80/20 train/test split (`seed=42`) was used, with a **3-fold CrossValidator** over a logistic-regression hyper-parameter grid (`regParam`, `elasticNetParam`) to document the tuning pipeline.

Algorithmic complexity (B1). Distributed reads/transforms are effectively $O(n / p)$ across p partitions; the shuffle-inducing `groupBy` / `join` stages are $O(n \log n)$ dominated by sort/exchange. Tree ensembles train in $O(t \cdot n \cdot f \cdot \log n)$ for t trees and f features; logistic regression is $O(i \cdot n \cdot f)$ for i iterations — consistent with its far shorter runtime.

6. System Design & Implementation

Architecture. A five-stage pipeline (Figure A) moves data from ingestion through processing, storage, modelling and serving.



Software stack. Apache Spark 3.5.3 (`local[*]`), Spark MLlib, Python 3.9, Java 17, SQLite, matplotlib/seaborn/plotly and Streamlit, under Git/GitHub version control. Configuration is centralised in `config/spark_config.py`.

Database. The SQLite star schema (Figure B) stores all 1,167,832 fact rows with indexed foreign keys, enabling efficient analytical joins.





Figure B. Relational star schema.

Big-data evidence. Spark was configured with **8 partitions** (> 4 required). The Storage tab (Figure C) confirms the cached DataFrame at **100 % fraction cached across 8 partitions (36.3 MiB)**; the Jobs tab shows 8/8 tasks per stage, evidencing balanced parallelism. Caching, `repartition`, broadcast joins and lazy DAG evaluation are all exercised in `spark_ui_demo.py`.



Security considerations. Every query is **parameterised** (? placeholders); an embedded test injects `x'`; `DROP TABLE fact_incident; --` and confirms zero rows match and the table remains intact — no injection surface. No credentials are hard-coded.

User interface. A Streamlit dashboard (Figure D) exposes KPIs, operational charts, an operator league table and the model-performance summary with an empirical risk lookup.




Figure D. Streamlit dashboard (operations view).

7. Results & Evaluation

Exploratory findings. Incidents peak sharply in the **07:00 AM** run with a secondary PM-dismissal peak (Figure 2); *Heavy Traffic* dominates reasons; and operator breakdown rates vary dramatically — from under 1 % to **over 60 %** — confirming strong predictive signal for benchmarking (Figure 3).




Figure 2. Incidents by hour of day, with peak bands shaded.




Figure 3. Operators ranked by breakdown rate (≥ 500 incidents).

Classification. All models were evaluated on the held-out test set (Table 1; Figure 4). The **GBTClassifier** performed best (ROC-AUC 0.847, F1 0.389), though all three achieved strong recall on the rare breakdown class (0.73–0.77) thanks to cost-sensitive weighting; precision is bounded by the 9.6 % base rate.

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
LogisticRegression	0.722	0.224	0.772	0.347	0.827	0.347
RandomForest	0.774	0.261	0.746	0.387	0.843	0.391
GBTClassifier	0.780	0.265	0.733	0.389	0.847	0.397

Table 1. Classification results (positive class = Breakdown).




Figure 4. ROC curves for the three classifiers.

Regression. The **Random-Forest regressor** achieved the lowest RMSE (**14.19 min**, $R^2 = 0.52$); GBT gave the best MAE (9.94 min) in one-third of the training time — a favourable **model-efficiency** trade-off (Table 2).

Model	RMSE	MAE	R^2	Train (s)
LinearRegression	16.42	11.93	0.359	4.8
RandomForestRegressor	14.19	10.04	0.521	72.4
GBTRegressor	14.34	9.94	0.511	21.6

Table 2. Regression results (delay minutes).

System evaluation. Stage runtimes were low (CSV load 1.1 s; cached count 2.3 s; broadcast join 0.6 s; Parquet write 5.9 s), confirming efficient use of the distributed engine. **Validation** against the data's own labels and the operator league table (cross-checked via independent SQL aggregation) corroborated the model rankings and feature-importance findings.

8. Critical Reflection

What worked well. The distributed pipeline handled 1.3 M rows comfortably; cost-sensitive learning produced usefully high breakdown recall; and the leakage-controlled feature design yielded an honest, realistic ROC-AUC rather than an artificially perfect score.

Challenges overcome. Messy free-text delays, corrupt timestamps, a Java-17 Spark-compatibility issue and websocket-based dashboard capture were each resolved methodically (documented in the repository).

Limitations — memory vs distributed. At ~1 M rows the data fits in memory, so Spark's per-task overhead makes it *slower* than Pandas would be for a single pass; Spark's value here is **scalability and reproducibility of pattern**, not raw speed at this size. The `on_time` proxy is coarse (the log records only incidents, not all trips), and the model cannot exceed the information in the recorded features.

Future work. Augmenting with weather and GTFS timetable feeds, threshold optimisation for the imbalanced classifier, and migration to a multi-node cluster for true horizontal scaling.

Ethical, legal & social considerations. The data is **CC0** and contains **no personal data** (only vehicle/route/operator codes), aligning with GDPR data-minimisation. Nonetheless, operator-level predictions risk **bias**: smaller operators with fewer records could be unfairly flagged, so outputs are framed as decision-support, not automated penalty. Transparency (open code, documented features) supports accountable use.

9. Conclusion

The project delivers a complete, reproducible big-data analytics platform that ingests over a million real records, engineers features at scale in PySpark, persists them securely to a relational star schema, and compares six machine-learning models to predict bus-service disruptions. Best-in-class results (ROC-AUC 0.847; RMSE 14.19 min) and the accompanying dashboard give a transport authority actionable insight into breakdown risk and delay behaviour. The work evidences all targeted learning outcomes — from algorithmic complexity and distributed programming to secure data handling and critical reflection — demonstrating that open transport data, processed responsibly at scale, is a practical asset for smarter, more reliable public transport.

10. References

He, H., & Garcia, E. A. (2009). Learning from imbalanced data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9), 1263–1284.

Moreira-Matias, L., Gama, J., Ferreira, M., Mendes-Moreira, J., & Damas, L. (2013). Predicting taxi-passenger demand using streaming data. *IEEE Transactions on Intelligent Transportation Systems*, 14(3), 1393–1402.

NYC Open Data. (2024). *Bus Breakdown and Delays* [Data set, id ez4e-fazm, CC0]. City of New York. <https://data.cityofnewyork.us/Transportation/Bus-Breakdown-and-Delays/ez4e-fazm>

Zaharia, M., Xin, R. S., Wendell, P., et al. (2016). Apache Spark: A unified engine for big data processing. *Communications of the ACM*, 59(11), 56–65.

11. Appendices

- **Appendix A — GitHub repository:** *[repository URL]* (all code, config, README, requirements, SQL dump, documentation).
- **Appendix B — Execution evidence:** docs/screenshots/exec_01–03_*.png (ingestion, ML, database + security console output); docs/screenshots/spark_ui_01–05_*.png (Spark UI Jobs, Stages, Storage, Executors, SQL).
- **Appendix C — Additional figures:** outputs/figures/fig01–fig15 (distributions, correlations, confusion matrix, feature importance, predicted-vs-actual, residuals).
- **Appendix D — Database:** database/schema.sql , database/sample_queries.sql , database/bus_reliability_sample_dump.sql .